

Image Classification

Object Detection for chess pieces

Linus Fackler and Fernando Colman

About the dataset

The dataset we are using is from <https://public.roboflow.com/object-detection/chess-full> (<https://public.roboflow.com/object-detection/chess-full>). The reason we chose this dataset was because it has the specific labeling type needed for training a YoloV7 model. The dataset consists of images of chessboard with chesspieces. There are 13 different classes: 'bishop', 'black-bishop', 'black-king', 'black-knight', 'black-pawn', 'black-queen', 'black-rook', 'white-bishop', 'white-king', 'white-knight', 'white-pawn', 'white-queen', 'white-rook'

What the model should be able to predict

The model should be able to recognize chess pieces and identify which type of chess piece it is (bishop, pawn, rook, ...)

Preprocessing of Data

There is no preprocessing that needs to be done here.

Split into train, test, val

This step is also not necessary with this dataset, as it already comes with a premade split. There are 202 train images, 58 validation and 29 test.

```
In [1]: import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

import os
```

Sequential Model

Reading in data

I've had problems with my original chess dataset for this part of the project. That's why I am using a tensorflow example dataset, so that I can at least get some results.

```
In [38]: import numpy as np
import os
import PIL
import PIL.Image
import tensorflow as tf
import tensorflow_datasets as tfds
import pathlib

from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [127]: import pathlib
dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_
images/flower_photos.tgz"
data_dir = tf.keras.utils.get_file('flower_photos', origin=dataset_url, untar=
True)
data_dir = pathlib.Path(data_dir)
```

```
In [128]: image_count = len(list(data_dir.glob('*/*.jpg')))
print("Number of images: ", image_count)
```

Number of images: 3670

Here we took 270 images out of the list and put it into test

```
In [130]: roses = list(data_dir.glob('roses/*'))
PIL.Image.open(str(roses[0]))
```

Out[130]:



```
In [41]: # %cd ..
# %ls
# train_url = "/content/drive/MyDrive/Colab Notebooks/Chess dataset/train"
# test_url = "/content/drive/MyDrive/Colab Notebooks/Chess dataset/test"
# val_url = "/content/drive/MyDrive/Colab Notebooks/Chess dataset/valid"

# train_ds = tf.keras.utils.image_dataset_from_directory(train_url, labels=None,
# validation_split=None, image_size=(200, 134), batch_size=16)
# test_ds = tf.keras.utils.image_dataset_from_directory(test_url, labels=None,
# validation_split=None, image_size=(200, 134), batch_size=16)
# val_ds = tf.keras.utils.image_dataset_from_directory(val_url, labels=None,
# validation_split=None, image_size=(200, 134), batch_size=16)
```

Found 202 files belonging to 1 classes.

Found 29 files belonging to 1 classes.

Found 58 files belonging to 1 classes.

Creating a dataset

```
In [100]: batch_size = 32
img_height = 180
img_width = 180

train_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.4,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.

Using 2202 files for training.

```
In [101]: val_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.

Using 734 files for validation.

```
In [107]: test_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=None,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.

```
In [61]: class_names = train_ds.class_names
    print("Classes: ", class_names)
```

Classes: ['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']

Visualizing the Data

```
In [62]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```

roses



dandelion



tulips



sunflowers



dandelion



roses




dandelion



roses



tulips



Creating a model

```
In [80]: num_classes = len(class_names)

model = tf.keras.models.Sequential([
    tf.keras.layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    tf.keras.layers.Conv2D(16, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(32, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(64, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(num_classes)
])
```

```
In [73]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
flatten_1 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 512)	401920
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 512)	262656
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 5)	2565

```
=====
Total params: 667,141
Trainable params: 667,141
Non-trainable params: 0
=====
```

```
In [83]: model.compile(optimizer='adam',
                      loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                      metrics=['accuracy'])
```

Training the model

```
In [84]: history = model.fit(train_ds, batch_size=32, epochs=10, validation_data=val_ds)
```

```
Epoch 1/10
92/92 [=====] - 16s 75ms/step - loss: 1.2986 - accuracy: 0.4285 - val_loss: 1.0989 - val_accuracy: 0.5518
Epoch 2/10
92/92 [=====] - 7s 70ms/step - loss: 0.9951 - accuracy: 0.6063 - val_loss: 0.9612 - val_accuracy: 0.6172
Epoch 3/10
92/92 [=====] - 7s 72ms/step - loss: 0.8066 - accuracy: 0.6928 - val_loss: 0.8758 - val_accuracy: 0.6635
Epoch 4/10
92/92 [=====] - 7s 73ms/step - loss: 0.5961 - accuracy: 0.7766 - val_loss: 0.8428 - val_accuracy: 0.6839
Epoch 5/10
92/92 [=====] - 7s 70ms/step - loss: 0.3783 - accuracy: 0.8719 - val_loss: 1.0283 - val_accuracy: 0.6730
Epoch 6/10
92/92 [=====] - 7s 71ms/step - loss: 0.2230 - accuracy: 0.9264 - val_loss: 1.2306 - val_accuracy: 0.6662
Epoch 7/10
92/92 [=====] - 7s 71ms/step - loss: 0.1255 - accuracy: 0.9642 - val_loss: 1.4477 - val_accuracy: 0.6471
Epoch 8/10
92/92 [=====] - 7s 71ms/step - loss: 0.0692 - accuracy: 0.9792 - val_loss: 1.5392 - val_accuracy: 0.6499
Epoch 9/10
92/92 [=====] - 7s 71ms/step - loss: 0.0488 - accuracy: 0.9854 - val_loss: 1.8818 - val_accuracy: 0.6158
Epoch 10/10
92/92 [=====] - 9s 100ms/step - loss: 0.0634 - accuracy: 0.9833 - val_loss: 1.7081 - val_accuracy: 0.6376
```


```
In [85]: history.history.keys()
```

```
Out[85]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

Evaluate the model

```
In [88]: import matplotlib.pyplot as plt

# Plot training & validation accuracy values
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



CNN architecture

We are going to use the same dataset as we did for the sequential.

```
In [97]: num_filters = 8
filter_size = 3
pool_size = 2

model2 = tf.keras.models.Sequential(
    [
        tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
        tf.keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(num_classes, activation="softmax"),
    ]
)
```


In [94]: `model2.summary()`

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_3 (MaxPooling 2D)	(None, 15, 15, 32)	0
conv2d_4 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_4 (MaxPooling 2D)	(None, 6, 6, 64)	0
flatten_3 (Flatten)	(None, 2304)	0
dropout_4 (Dropout)	(None, 2304)	0
dense_8 (Dense)	(None, 5)	11525
=====		
Total params: 30,917		
Trainable params: 30,917		
Non-trainable params: 0		

Train model

```
In [98]: model2.compile(optimizer='adam',
                        loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                        metrics=['accuracy'])

history2 = model2.fit(train_ds, batch_size=32, epochs=10, validation_data=val_ds)
```

Epoch 1/10

```
/usr/local/lib/python3.8/dist-packages/tensorflow/python/util/dispatch.py:108
2: UserWarning: "`sparse_categorical_crossentropy` received `from_logits=True`
`, but the `output` argument was produced by a sigmoid or softmax activation
and thus does not represent logits. Was this intended?"
    return dispatch_target(*args, **kwargs)
```

```
92/92 [=====] - 10s 93ms/step - loss: 55.3183 - accuracy: 0.2708 - val_loss: 1.6131 - val_accuracy: 0.2766
```

Epoch 2/10

```
92/92 [=====] - 8s 79ms/step - loss: 1.4799 - accuracy: 0.3784 - val_loss: 1.6985 - val_accuracy: 0.2779
```

Epoch 3/10

```
92/92 [=====] - 11s 112ms/step - loss: 1.2903 - accuracy: 0.4704 - val_loss: 1.8022 - val_accuracy: 0.3011
```

Epoch 4/10

```
92/92 [=====] - 7s 74ms/step - loss: 1.1900 - accuracy: 0.5228 - val_loss: 2.1110 - val_accuracy: 0.3052
```

Epoch 5/10

```
92/92 [=====] - 7s 75ms/step - loss: 1.0196 - accuracy: 0.6144 - val_loss: 2.2935 - val_accuracy: 0.3542
```

Epoch 6/10

```
92/92 [=====] - 7s 72ms/step - loss: 0.9315 - accuracy: 0.6458 - val_loss: 2.7418 - val_accuracy: 0.3338
```

Epoch 7/10

```
92/92 [=====] - 7s 72ms/step - loss: 0.8230 - accuracy: 0.7006 - val_loss: 2.5255 - val_accuracy: 0.3706
```

Epoch 8/10

```
92/92 [=====] - 7s 74ms/step - loss: 0.7458 - accuracy: 0.7371 - val_loss: 2.5079 - val_accuracy: 0.3896
```

Epoch 9/10

```
92/92 [=====] - 7s 72ms/step - loss: 0.6754 - accuracy: 0.7650 - val_loss: 3.3713 - val_accuracy: 0.3202
```


Epoch 10/10

```
92/92 [=====] - 7s 72ms/step - loss: 0.6966 - accuracy: 0.7619 - val_loss: 3.1006 - val_accuracy: 0.3529
```

Evaluate the model


```
In [99]: plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')
```

Out[99]: <matplotlib.legend.Legend at 0x7f06f028b610>



Pretrained Model

I've decided to use the YOLOv7 model for this project, as it seems to be the best open-source model for objects detection. I personally have experience with it, which was another factor as to why I chose this.



```
In [2]: # Downloading YOLO v7 Code
!git clone https://github.com/WongKinYiu/yolov7.git
%cd yolov7
!ls
```

Cloning into 'yolov7'...

remote: Enumerating objects: 1094, done.

remote: Counting objects: 100% (3/3), done.

remote: Compressing objects: 100% (3/3), done.

remote: Total 1094 (delta 0), reused 2 (delta 0), pack-reused 1091

Receiving objects: 100% (1094/1094), 69.85 MiB | 23.76 MiB/s, done.

Resolving deltas: 100% (521/521), done.

/content/yolov7

cfg	detect.py	hubconf.py	models	requirements.txt	tools	uti
-----	-----------	------------	--------	------------------	-------	-----

ls

data	export.py	inference	paper	scripts	train_aux.py
------	-----------	-----------	-------	---------	--------------

deploy	figure	LICENSE.md	README.md	test.py	train.py
--------	--------	------------	-----------	---------	----------

```
In [36]: # Pre-trained weights
!wget https://github.com/WongKinbYiu/yolov7/releases/download/v0.1/yolov7.pt

--2022-12-05 05:00:24-- https://github.com/WongKinbYiu/yolov7/releases/download/v0.1/yolov7.pt
Resolving github.com (github.com)... 140.82.121.4
Connecting to github.com (github.com)|140.82.121.4|:443... connected.
HTTP request sent, awaiting response... 404 Not Found
2022-12-05 05:00:24 ERROR 404: Not Found.
```

Resizing images

I decided to later comment out this section, as yolov7 does its own image resizing. Originally I wanted to do this so that it would not train with 2048px images, which would take too long.

I still kept it for future references, as I think the code is very helpful. Also, in case I ever want to use a different pretrained model, this might become helpful.

```
In [12]: # from PIL import Image

# directory = "/content/drive/MyDrive/Colab Notebooks/Chess/train/images"

# basewidth = 640

# for filename in os.listdir(directory):
#     f = os.path.join(directory, filename)
#     img = Image.open(f)
#     wpercent = (basewidth / float(img.size[0]))
#     hsize = int((float(img.size[1]) * float(wpercent)))
#     img = img.resize((basewidth, hsize), Image.ANTIALIAS)
#     img.save(f)
```

Training

This is where the **transfer learning** comes into play. We will now train the pretrained model with a new type of image. The images won't consist any of the 80 previous classes from the YOLOv7 (which actually is from the coco dataset), now we will only have our 13 classes, which include the chess figures.

```
In [133]: %cd "/content/drive/MyDrive/Colab Notebooks/yolov7"  
!python train.py --epochs 30 --weights yolov7.pt --batch-size 16 --data "/content/drive/MyDrive/Colab Notebooks/Chess/data.yaml" --name yolov7-custom --device 0 --exist-ok
```

/content/drive/MyDrive/Colab Notebooks/yolov7

YOLOR 🚀 2022-12-5 torch 1.12.1+cu113 CUDA:0 (Tesla T4, 15109.75MB)

Namespace(adam=False, artifact_alias='latest', batch_size=16, bbox_interval=-1, bucket='', cache_images=False, cfg='', data='/content/drive/MyDrive/Colab Notebooks/Chess/data.yaml', device='0', entity=None, epochs=30, evolve=False, exist_ok=True, freeze=[0], global_rank=-1, hyp='data/hyp.scratch.p5.yaml', image_weights=False, img_size=[640, 640], label_smoothing=0.0, linear_lr=False, local_rank=-1, multi_scale=False, name='yolov7-custom', noautoanchor=False, nosave=False, notest=False, project='runs/train', quad=False, rect=False, resume=False, save_dir='runs/train/yolov7-custom', save_period=-1, single_cls=False, sync_bn=False, total_batch_size=16, upload_dataset=False, v5_metric=False, weights='yolov7.pt', workers=8, world_size=1)

tensorboard: Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>

hyperparameters: lr0=0.01, lrf=0.1, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=0.05, cls=0.3, cls_pw=1.0, obj=0.7, obj_pw=1.0, iou_t=0.2, anchor_t=4.0, fl_gamma=0.0, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.2, scale=0.9, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, mosaic=1.0, mixup=0.15, copy_paste=0.0, paste_in=0.15, loss_ota=1

wandb: Install Weights & Biases for YOLOR logging with 'pip install wandb' (recommended)

Overriding model.yaml nc=80 with nc=13

	from	n	params	module	a
arguments					
0	-1	1	928	models.common.Conv	
[3, 32, 3, 1]					
1	-1	1	18560	models.common.Conv	
[32, 64, 3, 2]					
2	-1	1	36992	models.common.Conv	
[64, 64, 3, 1]					
3	-1	1	73984	models.common.Conv	
[64, 128, 3, 2]					
4	-1	1	8320	models.common.Conv	
[128, 64, 1, 1]					
5	-2	1	8320	models.common.Conv	
[128, 64, 1, 1]					
6	-1	1	36992	models.common.Conv	
[64, 64, 3, 1]					
7	-1	1	36992	models.common.Conv	
[64, 64, 3, 1]					
8	-1	1	36992	models.common.Conv	
[64, 64, 3, 1]					
9	-1	1	36992	models.common.Conv	
[64, 64, 3, 1]					
10 [-1, -3, -5, -6]	1		0	models.common.Concat	
[1]					
11	-1	1	66048	models.common.Conv	
[256, 256, 1, 1]					
12	-1	1	0	models.common.MP	
[]					
13	-1	1	33024	models.common.Conv	
[256, 128, 1, 1]					
14	-3	1	33024	models.common.Conv	
[256, 128, 1, 1]					

15		-1	1	147712	models.common.Conv
[128, 128, 3, 2]					
16	[-1, -3]	1		0	models.common.Concat
[1]					
17		-1	1	33024	models.common.Conv
[256, 128, 1, 1]					
18		-2	1	33024	models.common.Conv
[256, 128, 1, 1]					
19		-1	1	147712	models.common.Conv
[128, 128, 3, 1]					
20		-1	1	147712	models.common.Conv
[128, 128, 3, 1]					
21		-1	1	147712	models.common.Conv
[128, 128, 3, 1]					
22		-1	1	147712	models.common.Conv
[128, 128, 3, 1]					
23	[-1, -3, -5, -6]	1		0	models.common.Concat
[1]					
24		-1	1	263168	models.common.Conv
[512, 512, 1, 1]					
25		-1	1	0	models.common.MP
[]					
26		-1	1	131584	models.common.Conv
[512, 256, 1, 1]					
27		-3	1	131584	models.common.Conv
[512, 256, 1, 1]					
28		-1	1	590336	models.common.Conv
[256, 256, 3, 2]					
29	[-1, -3]	1		0	models.common.Concat
[1]					
30		-1	1	131584	models.common.Conv
[512, 256, 1, 1]					
31		-2	1	131584	models.common.Conv
[512, 256, 1, 1]					
32		-1	1	590336	models.common.Conv
[256, 256, 3, 1]					
33		-1	1	590336	models.common.Conv
[256, 256, 3, 1]					
34		-1	1	590336	models.common.Conv
[256, 256, 3, 1]					
35		-1	1	590336	models.common.Conv
[256, 256, 3, 1]					
36	[-1, -3, -5, -6]	1		0	models.common.Concat
[1]					
37		-1	1	1050624	models.common.Conv
[1024, 1024, 1, 1]					
38		-1	1	0	models.common.MP
[]					
39		-1	1	525312	models.common.Conv
[1024, 512, 1, 1]					
40		-3	1	525312	models.common.Conv
[1024, 512, 1, 1]					
41		-1	1	2360320	models.common.Conv
[512, 512, 3, 2]					
42	[-1, -3]	1		0	models.common.Concat
[1]					
43		-1	1	262656	models.common.Conv

[1024, 256, 1, 1]				
44	-2	1	262656	models.common.Conv
[1024, 256, 1, 1]				
45	-1	1	590336	models.common.Conv
[256, 256, 3, 1]				
46	-1	1	590336	models.common.Conv
[256, 256, 3, 1]				
47	-1	1	590336	models.common.Conv
[256, 256, 3, 1]				
48	-1	1	590336	models.common.Conv
[256, 256, 3, 1]				
49	[-1, -3, -5, -6]	1	0	models.common.Concat
[1]				
50	-1	1	1050624	models.common.Conv
[1024, 1024, 1, 1]				
51	-1	1	7609344	models.common.SPPCSPC
[1024, 512, 1]				
52	-1	1	131584	models.common.Conv
[512, 256, 1, 1]				
53	-1	1	0	torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']				
54	37	1	262656	models.common.Conv
[1024, 256, 1, 1]				
55	[-1, -2]	1	0	models.common.Concat
[1]				
56	-1	1	131584	models.common.Conv
[512, 256, 1, 1]				
57	-2	1	131584	models.common.Conv
[512, 256, 1, 1]				
58	-1	1	295168	models.common.Conv
[256, 128, 3, 1]				
59	-1	1	147712	models.common.Conv
[128, 128, 3, 1]				
60	-1	1	147712	models.common.Conv
[128, 128, 3, 1]				
61	-1	1	147712	models.common.Conv
[128, 128, 3, 1]				
62	[-1, -2, -3, -4, -5, -6]	1	0	models.common.Concat
[1]				
63	-1	1	262656	models.common.Conv
[1024, 256, 1, 1]				
64	-1	1	33024	models.common.Conv
[256, 128, 1, 1]				
65	-1	1	0	torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']				
66	24	1	65792	models.common.Conv
[512, 128, 1, 1]				
67	[-1, -2]	1	0	models.common.Concat
[1]				
68	-1	1	33024	models.common.Conv
[256, 128, 1, 1]				
69	-2	1	33024	models.common.Conv
[256, 128, 1, 1]				
70	-1	1	73856	models.common.Conv
[128, 64, 3, 1]				
71	-1	1	36992	models.common.Conv
[64, 64, 3, 1]				

72	-1 1	36992	models.common.Conv
[64, 64, 3, 1]			
73	-1 1	36992	models.common.Conv
[64, 64, 3, 1]			
74	[-1, -2, -3, -4, -5, -6] 1	0	models.common.Concat
[1]			
75	-1 1	65792	models.common.Conv
[512, 128, 1, 1]			
76	-1 1	0	models.common.MP
[]			
77	-1 1	16640	models.common.Conv
[128, 128, 1, 1]			
78	-3 1	16640	models.common.Conv
[128, 128, 1, 1]			
79	-1 1	147712	models.common.Conv
[128, 128, 3, 2]			
80	[-1, -3, 63] 1	0	models.common.Concat
[1]			
81	-1 1	131584	models.common.Conv
[512, 256, 1, 1]			
82	-2 1	131584	models.common.Conv
[512, 256, 1, 1]			
83	-1 1	295168	models.common.Conv
[256, 128, 3, 1]			
84	-1 1	147712	models.common.Conv
[128, 128, 3, 1]			
85	-1 1	147712	models.common.Conv
[128, 128, 3, 1]			
86	-1 1	147712	models.common.Conv
[128, 128, 3, 1]			
87	[-1, -2, -3, -4, -5, -6] 1	0	models.common.Concat
[1]			
88	-1 1	262656	models.common.Conv
[1024, 256, 1, 1]			
89	-1 1	0	models.common.MP
[]			
90	-1 1	66048	models.common.Conv
[256, 256, 1, 1]			
91	-3 1	66048	models.common.Conv
[256, 256, 1, 1]			
92	-1 1	590336	models.common.Conv
[256, 256, 3, 2]			
93	[-1, -3, 51] 1	0	models.common.Concat
[1]			
94	-1 1	525312	models.common.Conv
[1024, 512, 1, 1]			
95	-2 1	525312	models.common.Conv
[1024, 512, 1, 1]			
96	-1 1	1180160	models.common.Conv
[512, 256, 3, 1]			
97	-1 1	590336	models.common.Conv
[256, 256, 3, 1]			
98	-1 1	590336	models.common.Conv
[256, 256, 3, 1]			
99	-1 1	590336	models.common.Conv
[256, 256, 3, 1]			
100	[-1, -2, -3, -4, -5, -6] 1	0	models.common.Concat

	2/29	10.9G	0.07564	0.03249	0.03979	0.1479	150	
640:	100%	13/13	[00:45<00:00, 3.53s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.12it/s]				
		all	58	386	0.0255	0.155	0.0	
192	0.00588							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize								
	3/29	10.9G	0.06944	0.03325	0.03913	0.1418	162	
640:	100%	13/13	[00:56<00:00, 4.33s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.05it/s]				
		all	58	386	0.0647	0.218	0.0	
797	0.0289							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize								
	4/29	10.9G	0.06427	0.03112	0.03728	0.1327	281	
640:	100%	13/13	[00:53<00:00, 4.13s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.01it/s]				
		all	58	386	0.152	0.525	0.	
117	0.0503							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize								
	5/29	10.9G	0.05963	0.03064	0.03706	0.1273	192	
640:	100%	13/13	[00:48<00:00, 3.70s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.04it/s]				
		all	58	386	0.118	0.597	0.	
173	0.0735							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize								
	6/29	10.9G	0.05459	0.02981	0.03654	0.1209	68	
640:	100%	13/13	[00:49<00:00, 3.78s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:02<00:00, 1.04s/it]				
		all	58	386	0.0963	0.444	0.	
178	0.0784							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize								
	7/29	10.9G	0.05172	0.03089	0.03647	0.1191	199	
640:	100%	13/13	[00:51<00:00, 3.95s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:02<00:00, 1.01s/it]				
		all	58	386	0.0744	0.79	0.	
139	0.0546							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize								
	8/29	10.9G	0.04735	0.03059	0.03579	0.1137	214	
640:	100%	13/13	[00:50<00:00, 3.88s/it]					
		Class	Images	Labels	P	R	mAP	

@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.15it/s]				
137	0.0744		all	58	386	0.0841	0.565	0.
	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
640:	9/29	10.9G	0.04544	0.03295	0.03565	0.114	207	
	100%	13/13	[00:43<00:00,	3.38s/it]				
	Class	Images	Labels		P	R	mAP	
@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.17it/s]				
199	0.108		all	58	386	0.0976	0.723	0.
	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
640:	10/29	10.9G	0.04738	0.03229	0.03611	0.1158	173	
	100%	13/13	[00:47<00:00,	3.62s/it]				
	Class	Images	Labels		P	R	mAP	
@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.16it/s]				
254	0.148		all	58	386	0.285	0.397	0.
	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
640:	11/29	10.9G	0.04267	0.03198	0.0354	0.11	162	
	100%	13/13	[00:40<00:00,	3.13s/it]				
	Class	Images	Labels		P	R	mAP	
@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.26it/s]				
272	0.15		all	58	386	0.339	0.368	0.
	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
640:	12/29	10.9G	0.04107	0.03301	0.03517	0.1092	163	
	100%	13/13	[00:46<00:00,	3.55s/it]				
	Class	Images	Labels		P	R	mAP	
@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.09it/s]				
234	0.0715		all	58	386	0.264	0.429	0.
	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
640:	13/29	10.9G	0.05665	0.0265	0.03527	0.1184	223	
	100%	13/13	[00:41<00:00,	3.16s/it]				
	Class	Images	Labels		P	R	mAP	
@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.17it/s]				
202	0.0978		all	58	386	0.971	0.146	0.
	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
640:	14/29	10.9G	0.04885	0.02844	0.03588	0.1132	141	
	100%	13/13	[00:41<00:00,	3.21s/it]				
	Class	Images	Labels		P	R	mAP	
@.5	mAP@.5:.95: 100%	2/2	[00:01<00:00,	1.30it/s]				
0.28	0.148		all	58	386	0.546	0.275	

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
15/29	10.9G	0.0465	0.03197	0.03545	0.1139	195	
640: 100%	13/13 [00:46<00:00, 3.55s/it]						
Class	Images	Labels	P	R	mAP		
@.5 mAP@.5:.95: 100%	2/2 [00:01<00:00, 1.11it/s]						
all	58	386	0.684	0.271	0.		
234	0.0663						

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
16/29	10.9G	0.0444	0.03241	0.03472	0.1115	172	
640: 100%	13/13 [00:44<00:00, 3.45s/it]						
Class	Images	Labels	P	R	mAP		
@.5 mAP@.5:.95: 100%	2/2 [00:01<00:00, 1.10it/s]						
all	58	386	0.456	0.488	0.		
339	0.16						

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
17/29	10.9G	0.04124	0.03679	0.03414	0.1122	186	
640: 100%	13/13 [00:46<00:00, 3.58s/it]						
Class	Images	Labels	P	R	mAP		
@.5 mAP@.5:.95: 100%	2/2 [00:01<00:00, 1.17it/s]						
all	58	386	0.301	0.372	0.		
313	0.167						

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
18/29	10.9G	0.03927	0.03395	0.03396	0.1072	112	
640: 100%	13/13 [00:39<00:00, 3.05s/it]						
Class	Images	Labels	P	R	mAP		
@.5 mAP@.5:.95: 100%	2/2 [00:01<00:00, 1.05it/s]						
all	58	386	0.266	0.434	0.		
367	0.185						

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
19/29	10.9G	0.04229	0.02978	0.03377	0.1058	191	
640: 100%	13/13 [00:40<00:00, 3.12s/it]						
Class	Images	Labels	P	R	mAP		
@.5 mAP@.5:.95: 100%	2/2 [00:01<00:00, 1.20it/s]						
all	58	386	0.997	0.161	0.		
264	0.144						

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
20/29	10.9G	0.04299	0.02566	0.03367	0.1023	133	
640: 100%	13/13 [00:43<00:00, 3.32s/it]						
Class	Images	Labels	P	R	mAP		
@.5 mAP@.5:.95: 100%	2/2 [00:01<00:00, 1.17it/s]						
all	58	386	0.494	0.261	0.		
307	0.162						

Epoch	gpu_mem	box	obj	cls	total	labels	img_s
-------	---------	-----	-----	-----	-------	--------	-------

	21/29	10.9G	0.03875	0.03302	0.03315	0.1049	152	
640:	100%	13/13	[00:44<00:00, 3.44s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.23it/s]				
		all	58	386	0.228	0.608	0.	
359	0.201							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize	22/29	10.9G	0.03694	0.03152	0.03262	0.1011	223	
640:	100%	13/13	[00:45<00:00, 3.46s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.26it/s]				
		all	58	386	0.29	0.499	0.	
388	0.239							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize	23/29	10.9G	0.03453	0.0353	0.03211	0.1019	169	
640:	100%	13/13	[00:42<00:00, 3.26s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.09it/s]				
		all	58	386	0.26	0.496	0.	
397	0.228							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize	24/29	10.9G	0.03372	0.03461	0.0318	0.1001	208	
640:	100%	13/13	[00:45<00:00, 3.52s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.19it/s]				
		all	58	386	0.254	0.627		
0.41	0.272							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize	25/29	10.9G	0.03151	0.03497	0.03123	0.09772	94	
640:	100%	13/13	[00:44<00:00, 3.45s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.17it/s]				
		all	58	386	0.258	0.567		
0.4	0.24							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize	26/29	10.9G	0.03143	0.03531	0.03146	0.0982	161	
640:	100%	13/13	[00:53<00:00, 4.12s/it]					
		Class	Images	Labels	P	R	mAP	
@.5	mAP@.5:.95:	100%	2/2	[00:01<00:00, 1.25it/s]				
		all	58	386	0.304	0.532		
0.41	0.277							

	Epoch	gpu_mem	box	obj	cls	total	labels	img_s
ize	27/29	10.9G	0.02955	0.03537	0.03061	0.09552	274	
640:	100%	13/13	[00:42<00:00, 3.25s/it]					
		Class	Images	Labels	P	R	mAP	

```

@.5  mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.33it/s]
      all          58          386          0.608          0.434          0.
413      0.26

      Epoch    gpu_mem      box      obj      cls      total      labels  img_s
ize
      28/29     10.9G    0.03062    0.0377    0.03067    0.09899          272
640: 100% 13/13 [00:47<00:00, 3.64s/it]
      Class      Images      Labels          P          R          mAP
@.5  mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.19it/s]
      all          58          386          0.626          0.411          0.
418      0.273


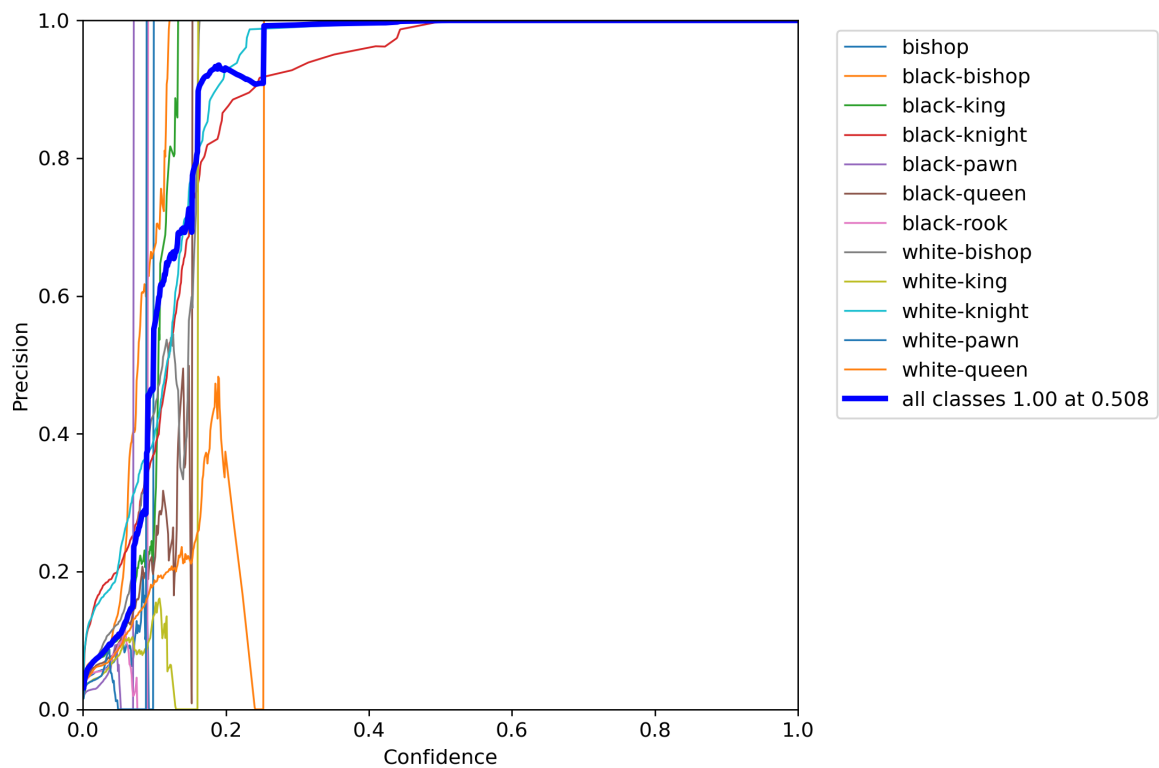
      Epoch    gpu_mem      box      obj      cls      total      labels  img_s
ize
      29/29     10.9G    0.02991    0.03217    0.03036    0.09244          141
640: 100% 13/13 [00:44<00:00, 3.43s/it]
      Class      Images      Labels          P          R          mAP
@.5  mAP@.5:.95: 100% 2/2 [00:03<00:00, 1.80s/it]
      all          58          386          0.617          0.441          0.
417      0.282
      black-bishop          58          22          1          0          0.
117      0.0723
      black-king          58          29          0.756          0.552          0.
754      0.574
      black-knight          58          30          0.653          0.3          0.
452      0.206
      black-pawn          58          77          0.436          1          0.
995      0.684
      black-queen          58          11          1          0          0.0
964      0.0748
      black-rook          58          28          0.283          0.25          0.
243      0.132
      white-bishop          58          22          1          0
0.12      0.089
      white-king          58          29          0.488          1          0.
669      0.49
      white-knight          58          19          0.145          0.32          0.
129      0.0729
      white-pawn          58          77          0.449          0.987          0.
985      0.729
      white-queen          58          16          1          0          0.
091      0.0638
      white-rook          58          26          0.194          0.885
0.35      0.193
30 epochs completed in 0.463 hours.

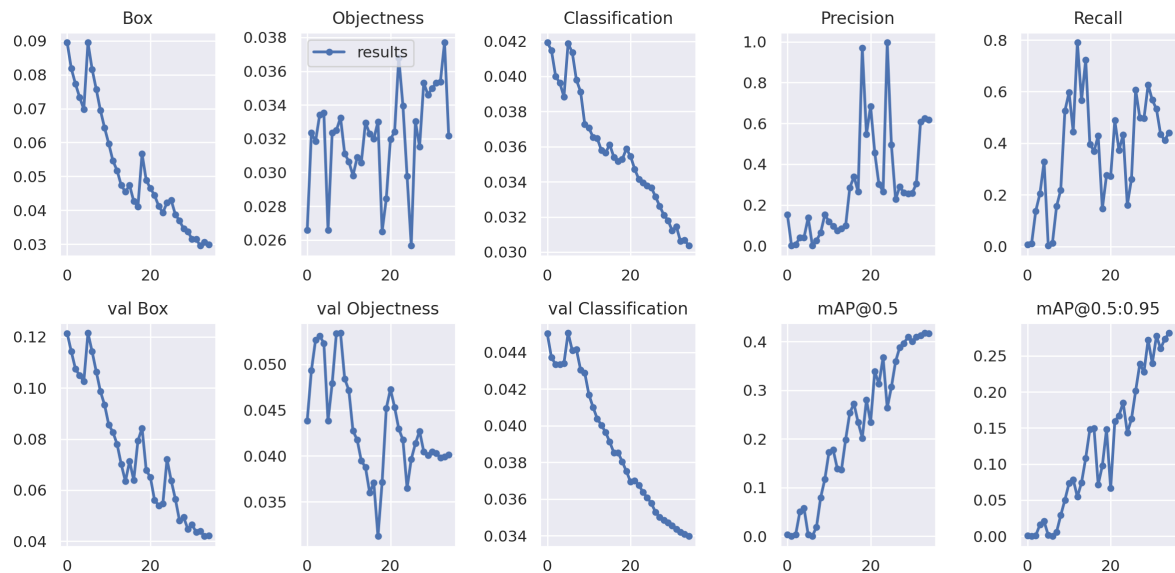
```

Optimizer stripped from runs/train/yolov7-custom/weights/last.pt, 74.9MB
Optimizer stripped from runs/train/yolov7-custom/weights/best.pt, 74.9MB


```
In [134]: from PIL import Image
pcurve = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/train/
yolov7-custom/P_curve.png')
confusionmatrix = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/ru
ns/train/yolov7-custom/confusion_matrix.png')
results = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/trai
n/yolov7-custom/results.png')

display(pcurve)
display(confusionmatrix)
display(results)
```





Testing

```
In [135]: python test.py --data "/content/drive/MyDrive/Colab Notebooks/Chess/data.yaml" --batch-size 16 --weights "/content/drive/MyDrive/Colab Notebooks/yolov7/runs/train/yolov7-custom/weights/best.pt" --exist-ok
```

```
Namespace(augment=False, batch_size=16, conf_thres=0.001, data='/content/drive/MyDrive/Colab Notebooks/Chess/data.yaml', device='', exist_ok=True, img_size=640, iou_thres=0.65, name='exp', no_trace=False, project='runs/test', save_conf=False, save_hybrid=False, save_json=False, save_txt=False, single_cls=False, task='val', v5_metric=False, verbose=False, weights=['/content/drive/MyDrive/Colab Notebooks/yolov7/runs/train/yolov7-custom/weights/best.pt'])
YOLOR 🚀 2022-12-5 torch 1.12.1+cu113 CUDA:0 (Tesla T4, 15109.75MB)
```

Fusing layers...

RepConv.fuse_repvgg_block

RepConv.fuse_repvgg_block

RepConv.fuse_repvgg_block

Model Summary: 306 layers, 36544546 parameters, 6194944 gradients

Convert model to Traced-model...

traced_script_module saved!

model is traced!

/usr/local/lib/python3.8/dist-packages/torch/functional.py:478: UserWarning: torch.meshgrid: in an upcoming release, it will be required to pass the indexing argument. (Triggered internally at ../aten/src/ATen/native/TensorShape.cpp:2894.)

return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]

val: Scanning '/content/drive/MyDrive/Colab Notebooks/Chess/valid/labels.cache' images and labels... 58 found, 0 missing, 0 empty, 0 corrupted: 100% 58/58 [00:00<?, ?it/s]



	Class	Images	Labels	P	R	mAP
@.5	mAP@.5:.95: 100%	4/4	[00:05<00:00, 1.41s/it]			
	all	58	386	0.612	0.439	0.
416	0.281					
	black-bishop	58	22	1	0	0.
117	0.072					
	black-king	58	29	0.75	0.552	
0.76	0.578					
	black-knight	58	30	0.651	0.3	0.
452	0.206					
	black-pawn	58	77	0.418	1	0.
995	0.684					
	black-queen	58	11	1	0	0.0
919	0.0707					
	black-rook	58	28	0.283	0.25	0.
244	0.135					
	white-bishop	58	22	1	0	
0.12	0.0891					
	white-king	58	29	0.488	1	0.
669	0.49					
	white-knight	58	19	0.136	0.298	0.
129	0.0713					
	white-pawn	58	77	0.429	0.987	0.
985	0.729					
	white-queen	58	16	1	0	0.0
807	0.0561					
	white-rook	58	26	0.194	0.885	0.
343	0.192					

Speed: 18.8/5.1/23.9 ms inference/NMS/total per 640x640 image at batch-size 16

Results saved to runs/test/exp

```
In [138]: pcurve = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/test/exp/P_curve.png')
confusionmatrix = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/test/exp/confusion_matrix.png')
#results = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/test/exp/results.png')

display(pcurve)
display(confusionmatrix)
#display(results)
```



The results here look terrible, *but* that was probably just because of the low number of epochs, which still took me forever to train. At first I tried it with 5 epochs, and it was even worse (I know, hard to imagine). With 300 epochs it would have probably been very good.

Detecting

Running sample images on the model on which we used transfer learning on.

```
In [139]: python detect.py --weights "/content/drive/MyDrive/Colab Notebooks/yolov7/runs/train/yolov7-custom/weights/best.pt" --conf 0.25 --img-size 640 --source "/content/drive/MyDrive/Colab Notebooks/Chess/test/images/a3863d0be6002c21b20ac88817b2c56f_jpg.rf.e421134b139d57e02e7df9468a35c1fb.jpg" --exist-ok
```

```
Namespace(agnostic_nms=False, augment=False, classes=None, conf_thres=0.25, device='', exist_ok=True, img_size=640, iou_thres=0.45, name='exp', no_trace=False, nosave=False, project='runs/detect', save_conf=False, save_txt=False, source='/content/drive/MyDrive/Colab Notebooks/Chess/test/images/a3863d0be6002c21b20ac88817b2c56f_jpg.rf.e421134b139d57e02e7df9468a35c1fb.jpg', update=False, view_img=False, weights=['/content/drive/MyDrive/Colab Notebooks/yolov7/runs/train/yolov7-custom/weights/best.pt'])
```

```
YOLOv7 2022-12-5 torch 1.12.1+cu113 CUDA:0 (Tesla T4, 15109.75MB)
```

```
Fusing layers...
```

```
RepConv.fuse_repvgg_block
```

```
RepConv.fuse_repvgg_block
```

```
RepConv.fuse_repvgg_block
```

```
Model Summary: 306 layers, 36544546 parameters, 6194944 gradients
```

```
Convert model to Traced-model...
```

```
traced_script_module saved!
```

```
model is traced!
```

```
/usr/local/lib/python3.8/dist-packages/torch/functional.py:478: UserWarning: torch.meshgrid: in an upcoming release, it will be required to pass the indexing argument. (Triggered internally at ../aten/src/ATen/native/TensorShape.cpp:2894.)
```

```
    return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
```

```
6 black-pawns, 5 white-pawns, Done. (17.9ms) Inference, (1.6ms) NMS
```

```
The image with the result is saved in: runs/detect/exp/a3863d0be6002c21b20ac88817b2c56f_jpg.rf.e421134b139d57e02e7df9468a35c1fb.jpg
```

```
Done. (0.323s)
```

```
In [141]: pil_im1 = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/test/exp/test_batch0_labels.jpg')
pil_im2 = Image.open('/content/drive/MyDrive/Colab Notebooks/yolov7/runs/test/exp/test_batch2_labels.jpg')
display(pil_im1)
display(pil_im2)
```

Output hidden; open in <https://colab.research.google.com> to view.

Analysis

Since I have had problems with my originally chosen dataset, I will not be able to compare the Sequential/RNN with the pretrained yolov7 model.

Generally, I can say, after working for quite a few time with a pretrained model, such as the YOLOv7, and before that the YOLOv5, I enjoy this more than building a model from scratch/using tensorflow's sequential models. Not only are pretrained models higher in accuracy (well, I couldn't prove it this time, but there's a reason as to why those are so popular), but also you will save time by just going with the (probably objectively) better choice. Especially, if you want to use features like opencv and detect objects in a real life setting with your camera or a video recording.

There is still the option to, as we have done it here, use transfer learning to use an existing model, like the YOLOv7, and "customize it" for your dataset. I have personally done this over the course of this semester, as I was part of the ACM Research. We used the FLIR thermal imaging dataset to make our own model (YOLOvCAPY), which works best for thermal images, especially when detecting objects while driving a car, which is what the dataset was made for.

For a general purpose like this here, the Sequential or RNN model seemed alright, but as soon as you want to expand, it seems impossible with these limited resources.